

CS 3430: SciComp with Py

Assignment 02

LU-Decomposition

Vladimir Kulyukin
Department of Computer Science
Utah State University

January 18, 2020

Learning Objectives

1. LU-Decomposition
2. Solving Linear Systems with LU-Decomposition

Introduction

In this assignment, we'll implement LU-Decomposition and use it to solve linear systems with the algorithms we discussed in Lectures 03 and 04. This assignment will give you more exposure to `numpy`. You'll save your coding solutions in `cs3430_s20_hw02.py` included in the zip and submit it in Canvas.

Problem 1: LU-Decomposition (1 point)

Implement the function `lu_decomp(A, n)` that does LU-Decomposition of an $n \times n$ matrix A and returns two $n \times n$ matrices U and L such that $LU = A$. For space efficiency, you can implement this function in such a way that A is destructively modified as it is being converted into U . Here's an example.

```
>>> a = np.array([[2, 3, -1],
... [0, 1, -3],
... [4, 5, -2]],
... dtype=float)
>>> u, l = lu_decomp(a, 3)
>>> u
array([[ 2.,  3., -1.],
       [ 0.,  1., -3.],
       [ 0.,  0., -3.]])
>>> l
array([[ 1.,  0.,  0.],
       [ 0.,  1.,  0.],
       [ 2., -1.,  1.]])
>>> np.dot(l, u)
array([[ 2.,  3., -1.],
       [ 0.,  1., -3.],
       [ 4.,  5., -2.]])
```

As we discussed in Lecture 04, we can implement two functions to streamline unit testing `lu_decomp(a, n)`. Let's define the function `comp_2d_mats(a, b, err=0.0001)` that compares two matrices `a` and `b` and returns `True` if `a` and `b` are of the same shape and, for every legitimate position `(i, j)`, `abs(a[i][j] - b[i][j]) <= err`, where `err` is a given level of error.

```
def comp_2d_mats(a, b, err=0.0001):
    ra, ca = a.shape
    rb, cb = b.shape
    if ra != rb:
        return False
    if ca != cb:
        return False
    for r in range(ra):
        for c in range(ca):
            if abs(a[r][c] - b[r][c]) > err:
                return False
    return True
```

We can now define the function `test_lud(a, err=0.0001, prnt_flag=True)` that takes a matrix `a`, checks that it is a square matrix, does LU-Decomposition of `a` with `lu_decomp()`, computes the *LU* product (i.e., the product of the two matrices returned by `lu_decomp()`), and uses `comp_2d_mats()` to compare the original matrix `a` with the *LU* product. If the printing flag `prnt_flag` is set to `True`, then matrices *U*, *L*, *LU*, and the original matrix are printed.

```
def test_lud(a, err=0.0001, prnt_flag=True):
    r, c = a.shape
    assert r == c
    u, l = lu_decomp(a.copy(), r)
    m2 = np.matmul(l, u)
    assert comp_2d_mats(a, m2)
    if prnt_flag:
        print('U:')
        print(u)
        print('L:')
        print(l)
        print('Original Matrix:')
        print(a)
        print('L*U:')
        print(m2)
```

Below are a couple of unit tests.

```
>>> a = np.array([[2, 3, -1],
... [0, 1, -3],
... [4, 5, -2]],
... dtype=float)
>>> test_lud(a, err=0.0001)
U:
[[ 2.  3. -1.]
 [ 0.  1. -3.]
 [ 0.  0. -3.]]
```

```

L:
[[ 1.  0.  0.]
 [ 0.  1.  0.]
 [ 2. -1.  1.]]
Original Matrix:
[[ 2.  3. -1.]
 [ 0.  1. -3.]
 [ 4.  5. -2.]]
L*U:
[[ 2.  3. -1.]
 [ 0.  1. -3.]
 [ 4.  5. -2.]]

>>> a = np.array([[73, 136, 173, 112],
... [61, 165, 146, 14],
... [137, 43, 183, 73],
... [196, 40, 144, 31]],
... dtype=float)
>>> test_lud(a, err=0.0001)
U:
[[ 73.          136.          173.          112.         ]
 [  0.          51.35616438    1.43835616 -79.5890411 ]
 [  0.           0.        -135.72712723 -466.09895972]
 [  0.           0.           0.         295.71529613]]
L:
[[ 1.          0.          0.          0.         ]
 [ 0.83561644  1.          0.          0.         ]
 [ 1.87671233 -4.13256868  1.          0.         ]
 [ 2.68493151 -6.33128834  2.29420978  1.         ]]
Original Matrix:
[[ 73. 136. 173. 112.]
 [ 61. 165. 146.  14.]
 [137.  43. 183.  73.]
 [196.  40. 144.  31.]]
L*U:
[[ 73. 136. 173. 112.]
 [ 61. 165. 146.  14.]
 [137.  43. 183.  73.]
 [196.  40. 144.  31.]]

```

Problem 2: Solving Linear Systems with LU-Decomposition (2 points)

Implement the function `bsubst(a, n, b, m)` that uses back substitution to solve $ax = b_1, b_2, \dots, b_m$, where a is an $n \times n$ upper-triangular matrix, n is its dimension, and b is an $n \times m$ matrix of m $n \times 1$ vectors b_1, b_2, \dots, b_m . This function returns the $n \times m$ matrix x of m $n \times 1$ vectors x_1, x_2, \dots, x_m such that $ax_1 = b_1, ax_2 = b_2, \dots, ax_m = b_m$. Here's an example.

```

a = np.array([[1, 3, -1],
              [0, 2,  6],
              [0, 0, -15]],
              dtype=float)
b = np.array([[-4,  5],
              [10, 11],

```

```

        [-30, 28]],
        dtype=float)
>>> x = bsubst(a, 3, b, 2)
>>> x
array([[ 1.          , -30.16666667],
       [-1.          ,  11.1         ],
       [ 2.          , -1.86666667]])
>>> np.dot(a, x[:,0])
array([-4.,  10., -30.])
>>> np.dot(a, x[:,1])
array([ 5., 11., 28.])

```

Implement the function `fbsubst(a, n, b, m)` that uses forward substitution to solve $ax = b_1, b_2, \dots, b_m$, where a is an $n \times n$ lower-triangular matrix, n is its dimension, and b is an $n \times m$ matrix of m $n \times 1$ vectors b_1, b_2, \dots, b_m . This function returns the $n \times m$ matrix x of m $n \times 1$ vectors x_1, x_2, \dots, x_m such that $ax_1 = b_1, ax_2 = b_2, \dots, ax_m = b_m$. Here's an example.

```

a = np.array([[1, 0, 0],
              [2, 1, 0],
              [-1, 3, 1]],
              dtype=float)
b = np.array([[-4, 10],
              [ 2, 12],
              [ 4, 21]],
              dtype=float)
>>> x = fbsubst(a, 3, b, 2)
[[ -4.  10.]
 [ 10.  -8.]
 [-30.  55.]]
>>> np.dot(a, x[:,0])
array([-4.,  2.,  4.])
>>> np.dot(a, x[:,1])
array([10., 12., 21.])

```

Implement the function `lud_solve(a, n, b, m)` that applies Algorithm 1 we discussed in Lecture 04 to solve the linear system $ax = b_1, b_2, \dots, b_m$, where a is an $n \times n$ matrix, b is an $n \times m$ matrix of m $n \times 1$ vectors b_1, b_2, \dots, b_m . This function uses LU-Decomposition to factor a into U and L . Then it uses forward substitution to solve $Ly = b$ for y , uses back substitution to solve $Ux = y$ for x , and returns x , which is an $n \times m$ matrix of m $n \times 1$ vectors x_i such that $ax_1 = b_1, ax_2 = b_2, \dots, ax_m = b_m$.

To test this function, we can define the function `check_lin_sys_sol()` that checks if a specific solution really solves the linear system at a given error level, as we discussed in class. The first four parameters in this function are the same as in `lud_solve`, x is an $n \times m$ matrix of m $n \times 1$ vectors x_i such that $ax_1 = b_1, ax_2 = b_2, \dots, ax_m = b_m$, and `err` is a given error level.

```

def check_lin_sys_sol(a, n, b, m, x, err=0.0001):
    ra, ca = a.shape
    assert ra == n
    assert ca == n
    assert b.shape[0] == n
    assert b.shape[1] == m
    assert b.shape == x.shape

```

```

for c in range(m):
    bb = np.array([np.matmul(a, x[:,c])]).T
    for r in range(n):
        assert abs(b[r][c] - bb[r][0]) <= err

```

We can use `check_lin_sys_sol` to define the function `test_lud_solve` to test `lud_solve`. The first four parameters are the same as in `check_lin_sys_sol`, the last two keyword arguments specify an accepted error level and a print flag.

```

def test_lud_solve(a, n, b, m, err=0.0001, prnt_flag=True):
    x = lud_solve(a, n, b, m)
    check_lin_sys_sol(a, n, b, m, x, err=err)
    if prnt_flag:
        print('A:')
        print(a)
        print('b:')
        print(b)
        print('x:')
        print(x)
        print('A*x:')
        print(np.dot(a, x))

```

Here are a couple of tests.

```

a = np.array([[1, 3, -1],
              [2, 8, 4],
              [-1, 3, 4]],
              dtype=float)
b = np.array([-4],
              [2],
              [4]],
              dtype=float)
>>> test_lud_solve(a, 3, b, 1)
A:
[[ 1.  3. -1.]
 [ 2.  8.  4.]
 [-1.  3.  4.]]
b:
[[-4.]
 [ 2.]
 [ 4.]]
x:
[[ 1.]
 [-1.]
 [ 2.]]
A*x:
[[-4.]
 [ 2.]
 [ 4.]]

a = np.array([[ 73., 136., 173., 112.],
              [ 61., 165., 146., 14.],
              [137., 43., 183., 73.]])

```

```

        [196., 40., 144., 31.]])
b = np.array([[4.0, 1.0],
              [-1.0, 2.0],
              [3.0, 3.0],
              [5.0, 4.0]])
>>> test_lud_solve(a, 4, b, 2)
A:
[[ 73. 136. 173. 112.]
 [ 61. 165. 146. 14.]
 [137. 43. 183. 73.]
 [196. 40. 144. 31.]]
b:
[[ 4. 1.]
 [-1. 2.]
 [ 3. 3.]
 [ 5. 4.]]
x:
[[ 0.04757985  0.01383696]
 [ 0.01254129 -0.00353939]
 [-0.04682867  0.01351588]
 [ 0.06180728 -0.01666954]]
A*x:
[[ 4. 1.]
 [-1. 2.]
 [ 3. 3.]
 [ 5. 4.]]

```

Implement the function `lud_solve2(u, l, n, b, m)` that applies Algorithm 2 we discussed in Lecture 04 to solve the linear system $ax = b_1, b_2, \dots, b_m$, where a is an $n \times n$ matrix, b is an $n \times m$ matrix of m $n \times 1$ vectors b_1, b_2, \dots, b_m . This function takes the U and L matrices obtained from the LU-Decomposition of a (these matrices are given in the first two parameters). The function uses L to convert b to c as would occur in the Gauss reduction of $[A|b]$ to $[U|c]$, uses back substitution to solve $Ux = c$ for x , and returns x , which is an $n \times m$ matrix of m $n \times 1$ vectors x_i such that $ax_1 = b_1, ax_2 = b_2, \dots, ax_m = b_m$.

We can define the function `test_lud_solve2` to test our implementation of `lud_solve2`. The procedure applies `lu_decomp` to factor a given matrix `a` into `u` and `l`, verifies the correctness of the solution at a given error level with `check_lin_sys_sol`, and, if the print flag is set to `True`, prints everything out.

```

def test_lud_solve2(a, n, b, m, err=0.0001, prnt_flag=True):
    aa = a.copy()
    u, l = lu_decomp(aa, n)
    bb = b.copy()
    x = lud_solve2(u, l, n, bb, m)
    check_lin_sys_sol(a, n, b, m, x, err=err)
    if prnt_flag:
        print('A:')
        print(a)
        print('b:')
        print(b)
        print('x:')
        print(x)

```

```

print('A*x:')
print(np.dot(a, x))

```

Here are a couple of tests.

```

a = np.array([[ 73., 136., 173., 112.],
               [ 61., 165., 146.,  14.],
               [137.,  43., 183.,  73.],
               [196.,  40., 144.,  31.]])
b = np.array([[4.0],
               [-1.0],
               [3.0],
               [5.0]])

```

```
>>> test_lud_solve2(a, 4, b, 1, err=err)
```

```

A:
[[ 73. 136. 173. 112.]
 [ 61. 165. 146.  14.]
 [137.  43. 183.  73.]
 [196.  40. 144.  31.]]

```

```

b:
[[ 4.]
 [-1.]
 [ 3.]
 [ 5.]]

```

```

x:
[[ 0.04757985]
 [ 0.01254129]
 [-0.04682867]
 [ 0.06180728]]

```

```

A*x:
[[ 4.]
 [-1.]
 [ 3.]
 [ 5.]]

```

```

a = np.array([[ 73., 136., 173., 112.],
               [ 61., 165., 146.,  14.],
               [137.,  43., 183.,  73.],
               [196.,  40., 144.,  31.]])
b = np.array([[4.0,  1.0],
               [-1.0, 2.0],
               [3.0,  3.0],
               [5.0,  4.0]])

```

```
>>> test_lud_solve2(a, 4, b, 2, err=err)
```

```

A:
[[ 73. 136. 173. 112.]
 [ 61. 165. 146.  14.]
 [137.  43. 183.  73.]
 [196.  40. 144.  31.]]

```

```

b:
[[ 4.  1.]
 [-1.  2.]
 [ 3.  3.]]

```

```

[ 5.  4.]]
x:
[[ 0.04757985  0.01383696]
 [ 0.01254129 -0.00353939]
 [-0.04682867  0.01351588]
 [ 0.06180728 -0.01666954]]
A*x:
[[ 4.  1.]
 [-1.  2.]
 [ 3.  3.]
 [ 5.  4.]]

```

Unit Testing

We're now in the position to test `lud_solve` on many linear systems. The pickled files `ab_5x5.pck`, `ab_50x50.pck`, and `ab_100x100.pck` contain 100 randomly generated square linear systems $[A \mid b]$ each. The name of the file specifies the size of the system. For example, in `ab_5x5.pck`, A is 5×5 and b is 5×1 . Let's define a function to load the pickled systems from a pck file.

```

import pickle
def load_lin_systems(file_name):
    with open(file_name, 'rb') as fp:
        return pickle.load(fp)

```

Here's how we can load all 100 linear systems from a given file.

```

>>> linsys = load_lin_systems('hw/hw02/ab_5x5.pck')
>>> len(linsys)
100
>>> A, b = linsys[0]
>>> A
array([[110., 176., 124.,  89., 193.],
       [162., 102.,  50., 125., 102.],
       [ 93., 117.,  66., 110., 164.],
       [  3.,  83., 156.,  73., 183.],
       [ 32., 137.,  51., 158.,  38.]])
>>> b
array([[ 36.],
       [165.],
       [116.],
       [156.],
       [125.]])
>>> A1, b1 = linsys[1]
>>> A1
array([[ 18.,  47., 119.,  12.,  64.],
       [ 93., 134.,  71.,  10., 113.],
       [187.,  80., 152.,  92.,  75.],
       [ 11., 194.,  74., 120., 175.],
       [156., 147., 151., 122., 105.]])
>>> b1
array([[ 82.],
       [ 48.],

```



```
[174.],  
[168.],  
[173.]])
```

Let's define a function we can use to test `lud_solve` on these systems.

```
def test_lud_solve_on_lin_systems(file_name, err=0.0001):  
    print('Testing LUD on {} ...'.format(file_name))  
    lu_problems = []  
    lin_systems = load_lin_systems(file_name)  
    for A, b in lin_systems:  
        try:  
            test_lud_solve(A, A.shape[0], b, 1, err=err, prnt_flag=False)  
        except Exception as e:  
            print(e)  
            lu_problems.append((A, b))  
    print('{} LUD solve failures out of {}'.format(len(lu_problems), len(lin_systems)))
```

Below are the results we should see if everything is working correctly.

```
>>> test_lud_solve_on_lin_systems('hw/hw02/ab_5x5.pck')  
Testing LUD on hw/hw02/ab_5x5.pck ...  
0 LUD solve failures out of 100  
>>> test_lud_solve_on_lin_systems('hw/hw02/ab_50x50.pck')  
Testing LUD on hw/hw02/ab_50x50.pck ...  
0 LUD solve failures out of 100  
>>> test_lud_solve_on_lin_systems('hw/hw02/ab_100x100.pck')  
Testing LUD on hw/hw02/ab_100x100.pck ...  
0 LUD solve failures out of 100
```

What to Submit

Save all your code in `cs3430_s20_hw02.py` and submit it in Canvas.

Happy Hacking!